

Facial Expression Recognition with Keras

by Coursera Project Network

Completed by Suhaimi William Chan

Taught by: [Snehan Kekre](#), Machine Learning Instructor

About this Course

Build and train a convolutional neural network (CNN) in Keras from scratch to recognize facial expressions. The data consists of 48x48 pixel grayscale images of faces. The objective is to classify each face based on the emotion shown in the facial expression into one of seven categories (0=Angry, 1=Disgust, 2=Fear, 3=Happy, 4=Sad, 5=Surprise, 6=Neutral).

You will use OpenCV to automatically detect faces in images and draw bounding boxes around them. Once you have trained, saved, and exported the CNN, you will directly serve the trained model to a web interface and perform real-time facial expression recognition on video and image data. For this project, you'll get instant access to a cloud desktop with Python, Jupyter, and Keras pre-installed.

Course Objectives

In this course, we are going to focus on four learning objectives:

1. Develop a facial expression recognition model in Keras
2. Build and train a convolutional neural network (CNN)
3. Deploy the trained model to a web interface with Flask
4. Apply the model to real-time video streams and image data

By the end of this course, you will be able to build and train a convolutional neural network (CNN) in Keras from scratch to recognize facial expressions. Once you have trained, saved, and exported the CNN, you will directly serve the trained model to a web interface and perform real-time facial expression recognition on video and image data.

Course Structure

This course is divided into 3 parts:

1. Course Overview: This introductory reading material.
2. Facial Expression Recognition with Keras: This is the hands on project that we will work on in Rhyme.
3. Graded Quiz: This is the final assignment that you need to pass in order to finish the course successfully.

Project Structure

The hands on project on Facial Expression Recognition is divided into following tasks:

Task 1: Introduction and Overview

- Introduction to the data and overview of the project.
- See a demo of the final product you will build by the end of this project.
- Introduction to the Rhyme interface.
- Import essential modules and helper functions from [NumPy](#), [Matplotlib](#), and [Keras](#).

Task 2: Explore the Dataset

- Display some images from every expression type in the Emotion FER [dataset](#).
- Check for class imbalance problems in the training data.

Task 3: Generate Training and Validation Batches

- Generate batches of tensor image data with real-time data augmentation.
- Specify paths to training and validation image directories and generates batches of augmented data.

Task 4: Create a Convolutional Neural Network (CNN) Model

- Design a convolutional neural network with 4 convolution layers and 2 fully connected layers to predict 7 types of facial expressions.
- Use Adam as the optimizer, categorical crossentropy as the loss function, and accuracy as the evaluation metric.

Task 5: Train and Evaluate Model

- Train the CNN by invoking the **model.fit()** method.
- Use **ModelCheckpoint()** to save the weights associated with the higher validation accuracy.
- Observe live training loss and accuracy plots in Jupyter Notebook for Keras.

Task 6: Save and Serialize Model as JSON String

- Sometimes, you are only interested in the architecture of the model, and you don't need to save the weight values or the optimizer.
- Use **to_json()**, which uses a JSON string, to store the model architecture.

Task 7: Create a Flask App to Serve Predictions

- Use open-source code from "[Video Streaming with Flask Example](#)" to create a flask app to serve the model's prediction images directly to a web interface.

Task 8: Create a Class to Output Model Predictions

- Create a FacialExpressionModel class to load the model from the JSON file, load the trained weights into the model, and predict facial expressions.

Task 9: Design an HTML Template for the Flask App

- Design a basic template in HTML to create the layout for the Flask app.

Task 10: Use Model to Recognize Facial Expressions in Videos

- Run the **main.py** script to create the Flask app and serve the model's predictions to a web interface.
- Apply the model to saved videos on disk.

Completed Jupyter Notebook and Scripts

This .zip file contains all the code and data used in the project. It includes the FER 2013 dataset, completed Jupyter notebook for training, the Flask app to serve predictions, and other utility scripts. Please feel free to modify any and all aspects of the code to suit your needs.

[Project.zip](#)

Once you have downloaded and extracted Project.zip, make sure to install dependencies using pipenv with the provided Pipfile and execute all commands using pipenv. Also, please make sure to add the correct path to the video file in camera.py on line 11. Next, to install pipenv, the dependencies, and run the main.py file, execute the following commands from your terminal or command prompt, making sure to add the right paths where necessary:

- `cd \path\to\Project\`
- `pip install pipenv`
- `pipenv install`
- `pipenv run python3 main.py`



Files

Running

Clusters

Select items to perform actions on them.

☐ 0 /

<input type="checkbox"/>	templates
<input type="checkbox"/>	test
<input type="checkbox"/>	train
<input type="checkbox"/>	utils
<input type="checkbox"/>	videos
<input type="checkbox"/>	Facial_Expression_Training.ipynb
<input type="checkbox"/>	camera.py
<input type="checkbox"/>	haarcascade_frontalface_default.xml
<input type="checkbox"/>	main.py
<input type="checkbox"/>	model.png

Facial Expression Recognition with Keras

Task 1: Import Libraries

```
In [1]: import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
import utils
import os
%matplotlib inline

from tensorflow.keras.preprocessing.image import ImageDataGenerator
from tensorflow.keras.layers import Dense, Input, Dropout, Flatten, Conv2D
from tensorflow.keras.layers import BatchNormalization, Activation, MaxPooling2D
from tensorflow.keras.models import Model, Sequential
from tensorflow.keras.optimizers import Adam
from tensorflow.keras.callbacks import ModelCheckpoint, ReduceLROnPlateau
from tensorflow.keras.utils import plot_model


from IPython.display import SVG, Image
from livelossplot import PlotLossesTensorFlowKeras
import tensorflow as tf
print("Tensorflow version:", tf.__version__)
```

Tensorflow version: 2.1.0

Task 2: Plot Sample Image

```
In [3]: utils.datasets.fer.plot_example_images(plt).show()
```





```
In [5]: for expression in os.listdir("train/"):
        print(str(len(os.listdir("train/" + expression))) + " " + expression + " images")
```

```
3171 surprise images
7215 happy images
4965 neutral images
3995 angry images
4830 sad images
436 disgust images
4097 fear images
```

Task 3: Generate Training and Validation Batches

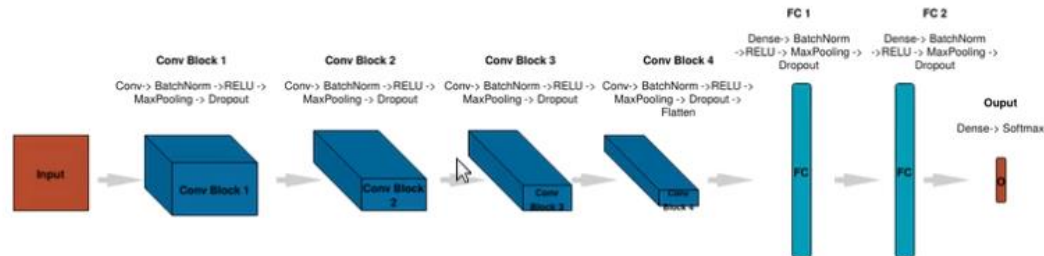
```
In [6]: img_size = 48
        batch_size = 64

        datagen_train = ImageDataGenerator(horizontal_flip=True)
        train_generator = datagen_train.flow_from_directory("train/",
                                                            target_size=(img_size, img_size),
                                                            color_mode='grayscale',
                                                            batch_size=batch_size,
                                                            class_mode='categorical',
                                                            shuffle=True)

        datagen_validation = ImageDataGenerator(horizontal_flip=True)
        validation_generator = datagen_train.flow_from_directory("test/",
                                                                target_size=(img_size, img_size),
                                                                color_mode='grayscale',
                                                                batch_size=batch_size,
                                                                class_mode='categorical',
                                                                shuffle=True)
```

```
Found 28709 images belonging to 7 classes.
Found 7178 images belonging to 7 classes.
```


Task 4: Create CNN Model



Inspired by Goodfellow, I.J., et.al. (2013). Challenged in representation learning: A report of three machine learning contests. *Neural Networks*, 64, 59-63. [doi:10.1016/j.neunet.2014.09.005](https://doi.org/10.1016/j.neunet.2014.09.005)

Task 4: Create CNN Model

```

```

Inspired by Goodfellow, I.J., et.al. (2013). Challenged in representation learning: A report of three machine learning contests. *Neural Networks*, 64, 59-63.

[\[doi:10.1016/j.neunet.2014.09.005\]](https://doi.org/10.1016/j.neunet.2014.09.005) (<https://arxiv.org/pdf/1307.0414.pdf>)

```
In [9]: model = Sequential()

# 1 - conv
model.add(Conv2D(64, (3,3), padding='same', input_shape=(48,48,1)))
model.add(BatchNormalization())
model.add(Activation('relu'))
model.add(MaxPooling2D(pool_size=(2,2)))
model.add(Dropout(0.25))

# 2 - conv layer
model.add(Conv2D(128, (5,5), padding='same'))
model.add(BatchNormalization())
model.add(Activation('relu'))
model.add(MaxPooling2D(pool_size=(2,2)))
model.add(Dropout(0.25))

# 3 - conv layer
model.add(Conv2D(512, (3,3), padding='same'))
model.add(BatchNormalization())
model.add(Activation('relu'))
model.add(MaxPooling2D(pool_size=(2,2)))
model.add(Dropout(0.25))
```

```

# 4 - conv layer
model.add(Conv2D(512, (3,3), padding='same'))
model.add(BatchNormalization())
model.add(Activation('relu'))
model.add(MaxPooling2D(pool_size=(2,2)))
model.add(Dropout(0.25))

model.add(Flatten())

model.add(Dense(256))
model.add(BatchNormalization())
model.add(Activation('relu'))
model.add(Dropout(0.25))

model.add(Dense(512))
model.add(BatchNormalization())
model.add(Activation('relu'))
model.add(Dropout(0.25))

model.add(Dense(7, activation='softmax'))

opt = Adam(lr=0.0005)
model.compile(optimizer=opt, loss='categorical_crossentropy', metrics=['accuracy'])
model.summary()

```

Model: "sequential_1"

Layer (type)	Output Shape	Param #
=====		
conv2d_4 (Conv2D)	(None, 48, 48, 64)	640
batch_normalization_4 (Batch Normalization)	(None, 48, 48, 64)	256
activation_5 (Activation)	(None, 48, 48, 64)	0
max_pooling2d_5 (MaxPooling2D)	(None, 24, 24, 64)	0
dropout_4 (Dropout)	(None, 24, 24, 64)	0
conv2d_5 (Conv2D)	(None, 24, 24, 128)	204928
batch_normalization_5 (Batch Normalization)	(None, 24, 24, 128)	512
activation_6 (Activation)	(None, 24, 24, 128)	0
max_pooling2d_6 (MaxPooling2D)	(None, 12, 12, 128)	0
dropout_5 (Dropout)	(None, 12, 12, 128)	0

conv2d_6 (Conv2D)	(None, 12, 12, 512)	590336
batch_normalization_6 (Batch Normalization)	(None, 12, 12, 512)	2048
activation_7 (Activation)	(None, 12, 12, 512)	0
token=4a7fa964736fc6eda836c240f6afdd481a5d359156d2d359		
max_pooling2d_8 (MaxPooling2D)	(None, 3, 3, 512)	0
dropout_7 (Dropout)	(None, 3, 3, 512)	0
flatten_1 (Flatten)	(None, 4608)	0
dense_1 (Dense)	(None, 256)	1179904
batch_normalization_8 (Batch Normalization)	(None, 256)	1024
activation_9 (Activation)	(None, 256)	0
dropout_8 (Dropout)	(None, 256)	0
dense_2 (Dense)	(None, 512)	131584
batch_normalization_9 (Batch Normalization)	(None, 512)	2048
activation_10 (Activation)	(None, 512)	0
dropout_9 (Dropout)	(None, 512)	0
dense_3 (Dense)	(None, 7)	3591
=====		
Total params: 4,478,727		
Trainable params: 4,474,759		
Non-trainable params: 3,968		


Task 6: Train and Evaluate Model

```
In [10]: epochs = 15
steps_per_epoch = train_generator.n//train_generator.batch_size
validation_steps = validation_generator.n//validation_generator.batch_size

checkpoint = ModelCheckpoint("model_weights.h5", monitor='val_accuracy',
                             save_weights_only=True, mode='max', verbose=1)
reduce_lr = ReduceLROnPlateau(monitor='val_loss', factor=0.1, patience=2,
                               min_lr=0.00001, model='auto')

callbacks = [PlotLossesTensorFlowKeras(), checkpoint, reduce_lr]

history = model.fit(
    x=train_generator,
    steps_per_epoch=steps_per_epoch,
    epochs=epochs,
    validation_data=validation_generator,
    validation_steps=validation_steps,
    callbacks=callbacks
)
```



Log-loss (cost function):

training (min: 0.864, max: 1.804, cur: 0.864)

validation (min: 0.961, max: 1.885, cur: 0.961)

accuracy:

training (min: 0.312, max: 0.677, cur: 0.677)

validation (min: 0.349, max: 0.643, cur: 0.643)

Epoch 00015: saving model to model_weights.h5

448/448 [=====] - 26s 59ms/step - loss: 0.8636 - accuracy: 0.6769 - val_loss: 0.9608 - val_accuracy: 0.6433

Task 7: Represent Model as JSON String

```
In [13]: model_json = model.to_json()
with open("model.json", "w") as json_file:
    json_file.write(model_json)
```

```

/home/rhyme/Desktop/Project/camera.py - Mousepad
File Edit Search View Document Help

import cv2
from model import FacialExpressionModel
import numpy as np

facec = cv2.CascadeClassifier('haarcascade_frontalface_default.xml')
model = FacialExpressionModel("model.json", "model_weights.h5")
font = cv2.FONT_HERSHEY_SIMPLEX

class VideoCamera(object):
    def __init__(self):
        self.video = cv2.VideoCapture("/home/rhyme/Desktop/Project/videos/facial_exp.mkv")

    def __del__(self):
        self.video.release()

    # returns camera frames along with bounding boxes and predictions
    def get_frame(self):
        _, fr = self.video.read()
        gray_fr = cv2.cvtColor(fr, cv2.COLOR_BGR2GRAY)
        faces = facec.detectMultiScale(gray_fr, 1.3, 5)

        for (x, y, w, h) in faces:
            fc = gray_fr[y:y+h, x:x+w]

            roi = cv2.resize(fc, (48, 48))
            pred = model.predict_emotion(roi[np.newaxis, :, :, np.newaxis])

            cv2.putText(fr, pred, (x, y), font, 1, (255, 255, 0), 2)
            cv2.rectangle(fr, (x,y), (x+w,y+h), (255,0,0), 2)

        _, jpeg = cv2.imencode('.jpg', fr)
        return jpeg.tobytes()

```


For your own live cam, you can change the following

`self.video = cv2.VideoCapture(0)`

```

* /home/rhyme/Desktop/Project/camera.py - Mousepad
File Edit Search View Document Help

import cv2
from model import FacialExpressionModel
import numpy as np

facec = cv2.CascadeClassifier('haarcascade_frontalface_default.xml')
model = FacialExpressionModel("model.json", "model_weights.h5")
font = cv2.FONT_HERSHEY_SIMPLEX

class VideoCamera(object):
    def __init__(self):
        self.video = cv2.VideoCapture(0)

    def __del__(self):
        self.video.release()

    # returns camera frames along with bounding boxes and predictions
    def get_frame(self):
        _, fr = self.video.read()
        gray_fr = cv2.cvtColor(fr, cv2.COLOR_BGR2GRAY)
        faces = facec.detectMultiScale(gray_fr, 1.3, 5)

        for (x, y, w, h) in faces:
            fc = gray_fr[y:y+h, x:x+w]

            roi = cv2.resize(fc, (48, 48))
            pred = model.predict_emotion(roi[np.newaxis, :, :, np.newaxis])

            cv2.putText(fr, pred, (x, y), font, 1, (255, 255, 0), 2)
            cv2.rectangle(fr, (x,y), (x+w,y+h), (255,0,0), 2)

        _, jpeg = cv2.imencode('.jpg', fr)
        return jpeg.tobytes()

```

```

/home/rhyme/Desktop/Project/main.py - Mousepad
File Edit Search View Document Help
from flask import Flask, render_template, Response
from camera import VideoCamera

app = Flask(__name__)

@app.route('/')
def index():
    return render_template('index.html')

def gen(camera):
    while True:
        frame = camera.get_frame()
        yield (b'--frame\r\n'
              b'Content-Type: image/jpeg\r\n\r\n' + frame + b'\r\n\r\n')

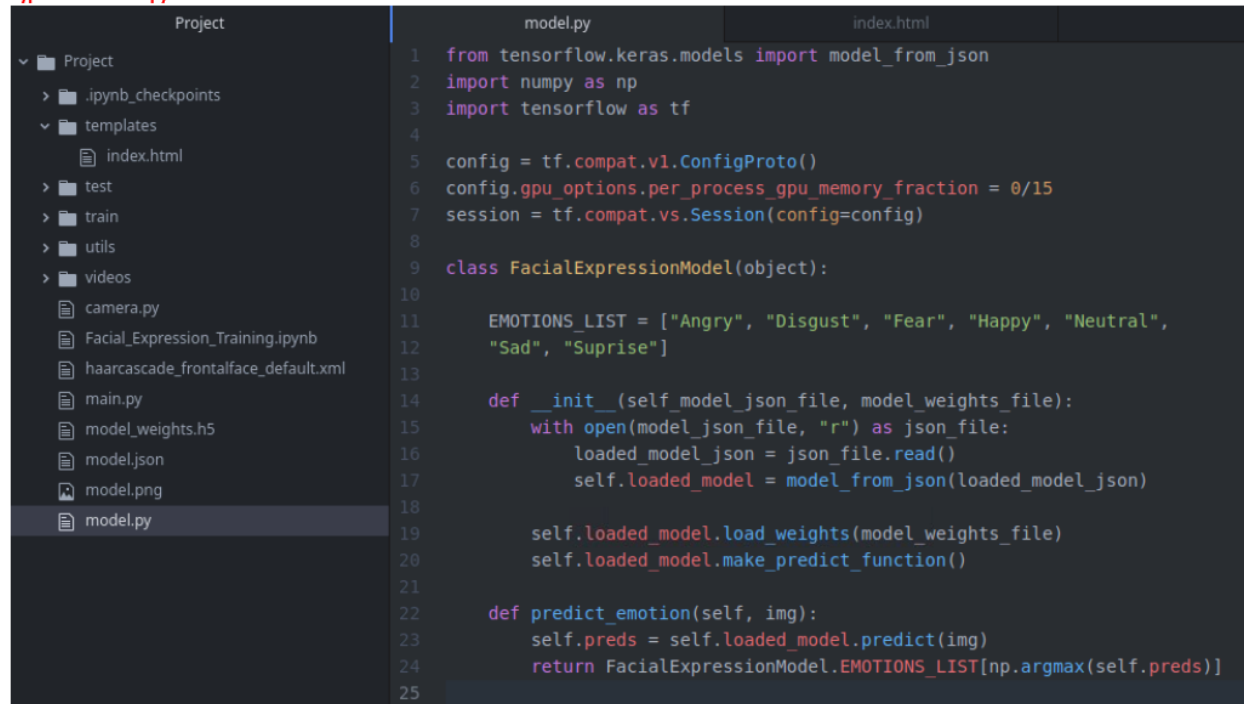
@app.route('/video_feed')
def video_feed():
    return Response(gen(VideoCamera()),
                    mimetype='multipart/x-mixed-replace; boundary=frame')

if __name__ == '__main__':
    app.run(host='0.0.0.0', debug=True)

```


Open terminal

Type: atom model.py



```
1 from tensorflow.keras.models import model_from_json
2 import numpy as np
3 import tensorflow as tf
4
5 config = tf.compat.v1.ConfigProto()
6 config.gpu_options.per_process_gpu_memory_fraction = 0/15
7 session = tf.compat.v1.Session(config=config)
8
9 class FacialExpressionModel(object):
10
11     EMOTIONS_LIST = ["Angry", "Disgust", "Fear", "Happy", "Neutral",
12                     "Sad", "Suprise"]
13
14     def __init__(self_model_json_file, model_weights_file):
15         with open(model_json_file, "r") as json_file:
16             loaded_model_json = json_file.read()
17             self.loaded_model = model_from_json(loaded_model_json)
18
19             self.loaded_model.load_weights(model_weights_file)
20             self.loaded_model.make_predict_function()
21
22     def predict_emotion(self, img):
23         self.preds = self.loaded_model.predict(img)
24         return FacialExpressionModel.EMOTIONS_LIST[np.argmax(self.preds)]
25
```

```

+ Enter the path for the new file.
templates/index.html

6 config.gpu_options.per_process_gpu_memory_fraction = 0/15
7 session = tf.compt.vs.Session(config=config)
8
9 class FacialExpressionModel(object):
10
11     EMOTIONS_LIST = ["Angry", "Disgust", "Fear", "Happy", "Neutral",
12                     "Sad", "Suprise"]
13
14     def __init__(self, model_json_file, model_weights_file):
15         with open(model_json_file, "r" as json_file:
16             loaded_model_json = json_file.read()
17             self.loaded = model_from_json(loaded_model_json)
18

```

```

Project
Project
> .ipynb_checkpoints
> templates
  index.html
> test
> train
> utils
> videos

model.py
1 <html>
2   <head>
3     <title>Facial Expression Recognition</title>
4   </head>
5   <body>
6     
7   </body>
8 </html>
9

```

Go back to Jupyter, go to Kernel and select restart
Then go back to terminal

```
Terminal - rhyme@ip-172-31-218-28: ~/Desktop/Project
File Edit View Terminal Tabs Help
rhyme@ip-172-31-218-28:~/Desktop/Project$ atom model.py
rhyme@ip-172-31-218-28:~/Desktop/Project$ cd..
cd..: command not found
rhyme@ip-172-31-218-28:~/Desktop/Project$ python3 main.py
```



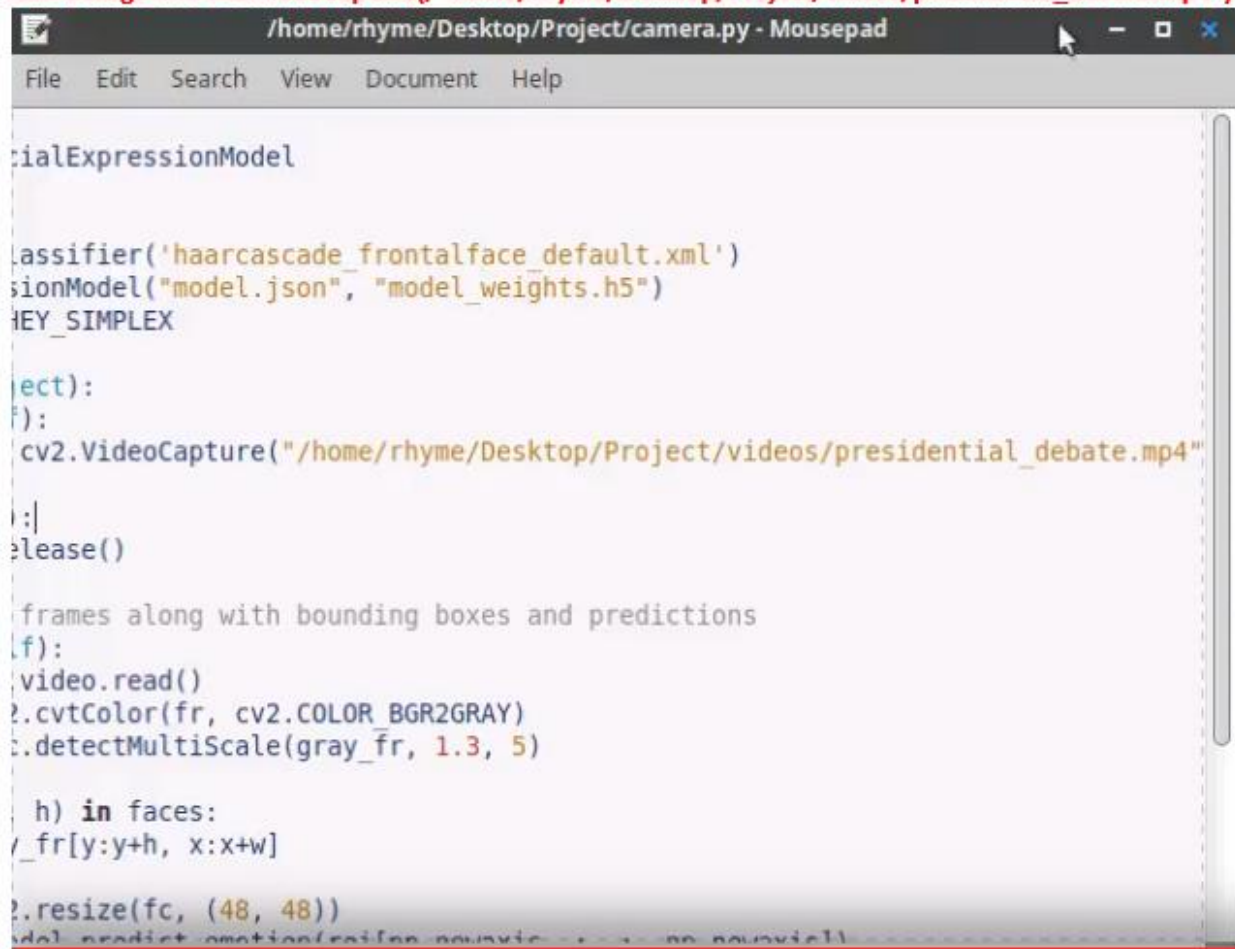






Change the location of file to be analyzed:

Just change it in `cv2.VideoCapture("/home/rhyme/Desktop/Project/videos/presidential_debate.mp4")`



```

FacialExpressionModel

    classifier('haarcascade_frontalface_default.xml')
    sionModel("model.json", "model_weights.h5")
    KEY_SIMPLEX

    (ect):
    f):
    cv2.VideoCapture("/home/rhyme/Desktop/Project/videos/presidential_debate.mp4")

    :|
    elease()

    frames along with bounding boxes and predictions
    f):
    video.read()
    2.cvtColor(fr, cv2.COLOR_BGR2GRAY)
    2.detectMultiScale(gray_fr, 1.3, 5)

    h) in faces:
    /_fr[y:y+h, x:x+w]

    2.resize(fc, (48, 48))
    del predict_emotion(emotion_prediction)
  
```

